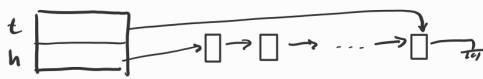


## Singly Linked List



```

p = f
if (p == null ptr) return p;
f = p->next;
if (f == null ptr) t = null ptr;
return p
  
```

- \* Reading followed by writing actions is what causes issues
- \* If you have instructions adding and removing then it gets worse

Can we use Atomic to make this work?

need some form of compare + swap  
↳ CMP XCHG  
(compare exchange)

```

p = f
if (p == null ptr) return p;
h = p->next;
if (f == null ptr) t = null ptr;
return p
  
```

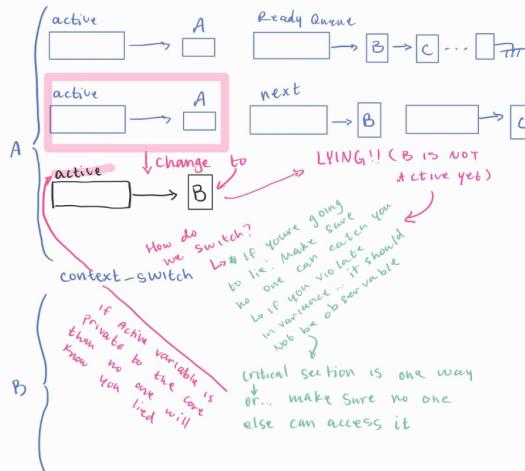
This is a point where the assertion might fail incorrectly.  
↳ This would be OK if everything ran sequentially.  
↳ BUT not in a parallel system.  
↳ Another thread might catch the invariant.

↳ need to use small critical section to avoid this issue

\* Assertions are saying "we can't be 100% sure that it works"

## Common Race Conditions

① Ex: Calling Yield() → Setting B to Active

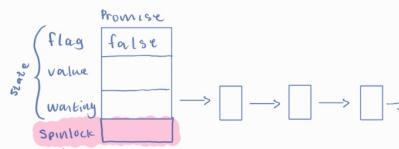


② Promise:



What happens if the value is set at the SAME time that someone is adding themselves to the waiting queue?

↳ no one will wake up the thread from the queue



Add a Spinlock! — hold the lock for the duration of the context switch

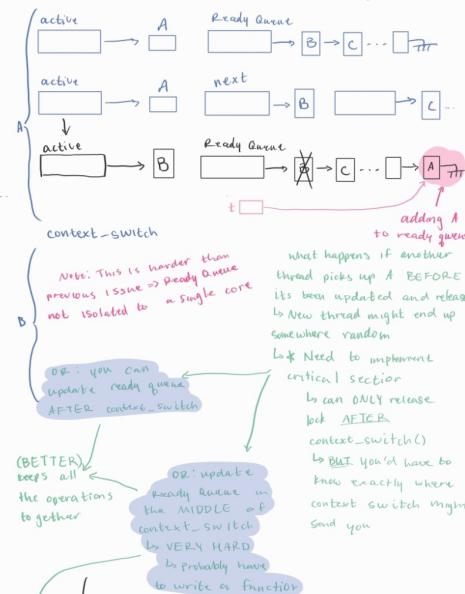
Double Checking Technique:

?? check the state prior to context-switch()  
if it's false then context-switch() and check after THEN decide what queue to be added

↳ you lock twice but for MUCH shorter time

\* in general: your lock should NOT depend on another lock  
↳ can produce a deadlock  
↳ no longer O(1)

③ Ex: calling Yield() → Adding A to ready queue



Note: This is harder than previous issue ⇒ ready queue not isolated to a single core  
↳ New thread might end up somewhere random  
↳ Need to implement critical section  
↳ can ONLY release lock AFTER context-switch()  
↳ BUT you'd have to know exactly where context switch might send you

Alternative Solution: Give each core an independent ready queue and distribute tasks to each core  
↳ high probability of distributing tasks unevenly  
↳ periodically rebalance tasks

How do you know how much work a task is?  
there's a global array where each core writes how much work a task is

↳ Race conditions don't matter

If there is gross imbalance then a busy core will put a portion of their tasks on a global queue  
↳ if it's out-of date its ok... nothing bad will happen

add the task before switch and use a flag to mark as dangerous

↳ during switch → change the flag

↳ less expensive than doing actual work during switch  
↳ if someone picks up task while dangerous it can just put it back  
↳ only a small waste

## Problem:

```
void notKernelMain() {
    Barrier b{N+3};
    for (i=0; i<N; i++) {
        thread[i] {
            print("before\n");
            b.sync();
            print("after\n");
        };
    };
    return;
}
```

very high probability these threads WONT finish before not kernel main() ends

This is a PROBLEM because b will get deleted: the usage of b

Need to capture b  $\Leftarrow$  outlives its existence but CANT use value

Makes copy of the barrier which is USELESS

### Solution:

```
void notKernelMain() {
    Barrier b{N+3};
    for (i=0; i<N; i++) {
        thread[i] {
            print("before\n");
            b.sync();
            print("after\n");
        };
    };
}
```

$b.sync()$  → makes sure all uses of b complete before ending notKernelMain()

## How do we make this Reliable?

```
void notKernelMain() {
    Barrier* b = new Barrier(N);
    for (i=0; i<N; i++) {
        thread[i] {
            print("before\n");
            b.sync();
            print("after\n");
        };
    };
    return;
}
```

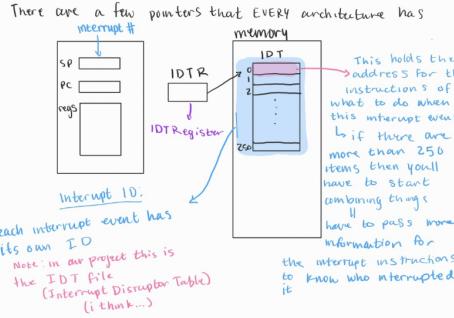
\* Put B in the heap (malloc)

\* creates a copy of the REFERENCE of b

\* NOTE: There is no reliable way to make sure b is deleted. ↳ would have to have extra sync and free OR implement that all threads (garbage) collector

## Ghetti lied to us in 429

\*Note: This will be discussing x86 architecture

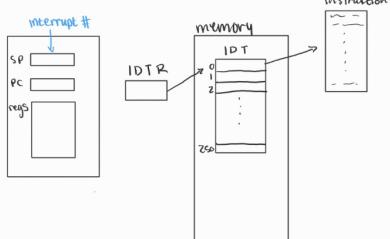


\* looks like a forced Function Call

Before taking an interrupt:

- saves 3 values to the current stack
  - Current PC
  - Flags Registers (ex. carry bit) → vulnerable: can be modified by one!!
  - CS register
- ↳ CS=Code Segment:
  - Describes the nature of the code that is running (ex. x64 or x32 mode)
  - ↳ Two interrupt might be running in a different mode

- x86 uses a special instruction (iret) to return
- interrupt handler will use push and pop to save EVERYTHING to the stack
  - ↳ interrupt return
  - ↳ push all / pop all



\* When an interrupt is called, the interrupt code is still running on the same hardware  
↳ This can cause a problem because it will continue to re-read the interrupt and loop infinitely

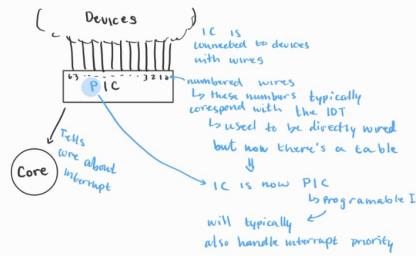
### Solution:

\* EFLAGS → x86 register

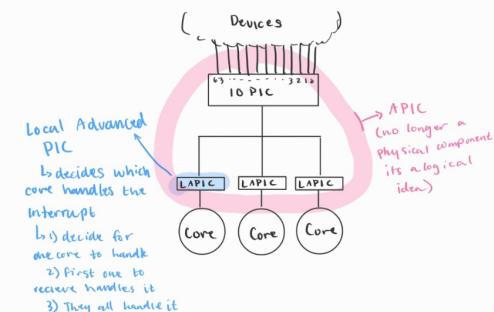
CLI → clear interrupt flag  
↳ interrupt flag tells you if you should ignore interrupts or not  
SLI → set interrupt flag  
↳ this is what defines if you're interrupted or not

\* There are maskable and non-maskable interrupts... ??

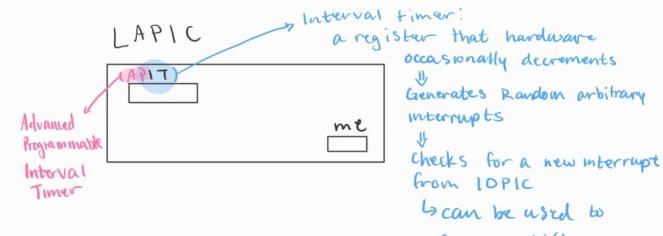
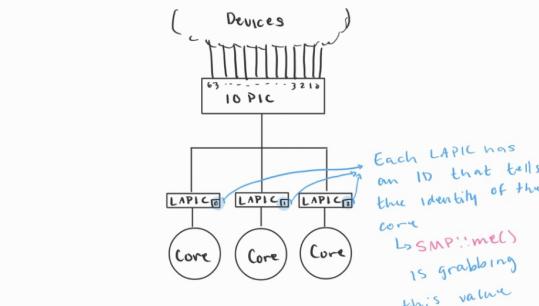
## Interrupt Controller



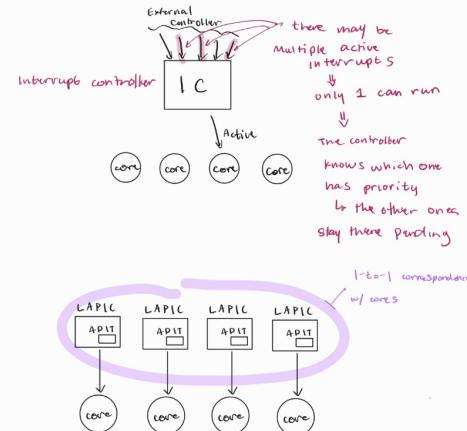
what if there are multiple cores



## APIC



Continued from last time...



- when there's an interrupt there's ATLEAST one active core
- There's an IF flag (interrupt flag) in each core that tells the core there is an interrupt
- interrupt won't be handled immediately but will happen "soon"

## HANDLING INTERRUPTS

- ① Turn off IF to prevent other interrupts from happening
- ② Execute the interrupt
- ③ iret is called → returns and pops things back off the stack (including the original IF flag)
  - ↳ you could manipulate the stack and change the IF flag
- \* The LAPIC will have a small buffer that stores a list of pending interrupts

Once the interrupt is done it sets the Flip-Flop in the LAPIC to allow the next interrupt through

EOI goes to IC to tell it that the interrupt has been handled and it can stop trying to get it handled

↳ register in the device

```
if (core == 0) { core so you don't count extra
    jiffies += 1; time
}
```

EOI

~~SEI()~~ ← DONT DO THIS!!  
re-enables IF which opens the possibility for multiple interrupts to be handled at one time

### Two Rules:

- ① Don't call ~~STI()~~
- ② Must run in O(1)

Reminder:  
CLI: disables Interrupts  
STI: enables

Note: you don't have to disable interrupts in program hardware will do for you

How to disable interrupt:

- \* must be O(1)
- \* spinlock must also be O(1)
- \* Aguarun boundary:
- yield()

↳ In our world:

no matter what if the timer reaches its limit → Force yield()  
↳ normally this would be optimized to not switch away from a task too quickly

Note: Test threads (threads created by test)  
should run w/ interrupts enabled

Problem: when you switch threads the thread may be assuming the state of IF but it may not have been restored  
↳ Need to store the des

```
APIT Handler:
{
    update jiffies
    EOI
    yield()
}
```

Question: How do we ensure IF is enabled?  
Solution: push IF flag during context\_switch

Problem: what if you interrupt in the middle of a yield()?

- \* will look like calling yield() recursively
- \* yield() is calling and manipulating global variables

what about if you're in the middle of spinlock()?  
↳ could result in deadlock

what about middle of malloc()?  
↳ could mess up malloc

CANT call yield() → yield NOT O(1)
 

- ↳ can you call yield() while yielding?  
↳ manipulates ready queue which has a spin lock which is NOT O(1)
- ↳ use "busy flag" to check if the thread is yielding  
↳ you don't know how long you will wait to obtain lock
- ↳ Only Fixes the yield problem
- \* one idea: have a local ready queue → allows you to use ready queue in O(1)

NOTE: This is ONLY for the Timer interrupt  
↳ can ignore the timer interrupt b/c you're already yielding

How does the thread know its identity?  
↳ TCB

```
auto my tcb = current_thread[SMP::me()]
    ↓ same as
auto id = SMP::me();
    ↑ [INTERRUPT]
    ↓
    auto my tcb = current_thread[id];
    ↓
    EASY SOLUTION
    CLI;
    auto id = SMP::me();
    auto my tcb = current_thread[id];
    STI();
```

BIG PROBLEM!  
SMP::me() changes!  
↓ RACE condition!

Disable interrupts

New Problem!

```
// how do you know interrupts were enabled before?
CLI;
auto id = SMP::me();
auto my tcb = current_thread[id];
STI();
```

wouldn't want to enable them here!  
you don't want to disable anywhere else  
\* Put this in a method instead of accessing current\_thread

Solution!

```
was = interrupts::disable();
    ↓
    Interrupts::restore(was);
```

disables and returns previous state  
\* This needs to be written

\* Note: Gheith recommends putting asserts everywhere to check the IF flag!!

NEW PROBLEM!

```
O(1)
    ↓ [Interrupt]
    ↓
    Spin.lock();
    :
    Spin.unlock();
```

↳ Can lead to deadlock

TWO BAD IMPLEMENTATIONS

```
Spin.lock();
disable();
Spin.lock();
enable();
Spin.unlock();
Spin.unlock();
```

can still interrupt here

```
disable();
Spin.lock();
Spin.unlock();
enable();
```

no longer O(1)

Solution

change spinlock implementation:

```
disable();
while(exchange(taken, 1)) {
    restore();
    // stuff
}
disable();
restore();
```

Disables interrupts for this line

\* Note: There may still be problems with this implementation  
Ex: what if a different thread unlocks?

\* Note: Relax the disable = O(1) requirement

How To Sleep?

Gheith's first implementation

```
void sleep(int n) {
    target = jiffies + n + jiffies;
    while(jiffies <= target) {
    }
}
```

estimated bell of time  
BAD!  
blocks the whole core from running

↓ timer Queue SORTED!

.check and update during every interrupt

.you technically could just have local timer queues cause they all being updated at the same time

↳ doesn't matter because still can't add everything to ready queue

\* A thread that calls sleep will always have interrupts enabled

↳ sleep will always be run by a user process (a test case)

\* anything you do in the APIT needs to be O(1)

↳ APIT can't create threads ⇒ not O(1)  
↳ give APIT a pre-made thread to handle the cleanup work that is slower

↳ this special thread should prevent preemption

this is ok  
↳ can't stop the interrupt but can be known  
↳ Stop the interrupt from doing anything it will yield  
↳ disabling interrupt handler stops time from advancing

\* This helper thread will add the things to the ready queue

\* helper thread ⇒ offline interrupt handler  
↳ will never go into ready queue  
↳ manually call it!

\* how do you communicate with helper thread?  
↳ put some sort of data structure in the thread to send it info  
↳ (e.g. timer ⇒ let the thread figure out who can wake up)

\* This is only needed if the next thread has interrupts disabled

↳ if interrupts are enabled in the next thread then it could just do the cleanup  
↳ if interrupts not enabled then it can't  
↳ you don't want interrupts to be disabled while cleaning up  
↳ is this too much work?