


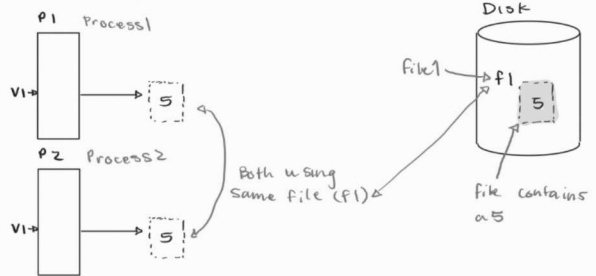
how do you mmap a particular address?

permissions (r, w, x) file descriptor
 mmap (address, size, flags, fd, offset)
 ↳ this is what makes it complicated -- there's a lot
 ↳ private / shared
 ↳ anonymous / file
 ↳ what's in fd and offset?
 ↳ system dependent
 ↳ some don't care and just ignore
 ↳ others need -1, etc.
 ↳ shared is much more complicated than private

Then each process can just point to the instance in cache



How do you synchronize file access?




*When another process mmmaps this file it will go to the cache and find it
 ↳ if exists --> point to it
 ↳ guaranteed to be updated
 ↳ else --> fetch from disk and put in cache
 ↳ slow !!
 *When we close the file...
 ↳ update it on disk...


[IDK]

*What happens when no one points to the file anymore?
 ↳ you could reference count and delete...
 ↳ but WHY?
 ↳ getting it from disk was so slow and the RAM memory space is already so big
 ↳ keep it in case someone else wants to mmap it later?

But what if P1 changes the file?




What about private mmap?



But... this will waste a lot of RAM
 not everyone will try and write
 so...

make a "temporary office" for everyone to share until someone tries to write
 *COW --> copy on write

How do we find the item in the cache?
 (what do we hash)

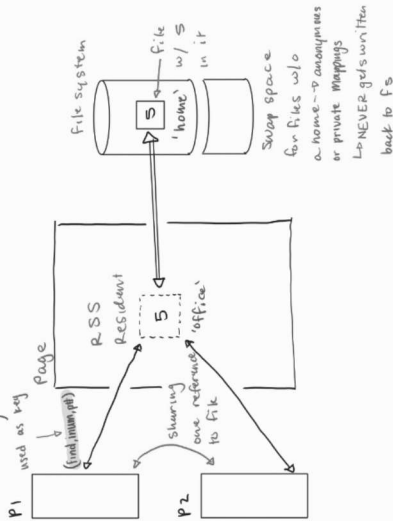


execel (path, 0)

- This means there are 0 args (argc == 0)
- This is valid - it's just weird cause there is typically at least one arg which is Smith

dynamic library : libc.so

static library : libc.a



What happens if you run out of physical frames?

- In p6 we would panic but... how what?

- write it back to FS and set the present bit in that process to 0
- ↳ next time it tries to access the file it will page fault
- we want to avoid clearing a page that will be used soon
- ↳ how do we know which file will be unlikely to be used?
- most optimal is to throw away the page you will need furthest in the future
- ↳ but... we don't know this...

- If the item in residence hasn't been written yet... you don't need to write back
- ↳ worthwhile optimization.

we need some sort of function to invert this...



converts the VPN to get a PPN



There's a similar hardware component that gets the VPNs pointing to a PPN

called IPT;

we will call it

PMAP-Physical Map

↳ Hashmap?

↳ there's a simpler way!

Array indexed by PPN

↳ each item points to a linked list

↳ or some other representation of a set

↳ set of physical memory

PMAP

PPN

PHYS MEM

kernel heap

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...

...


...

...

...


...

How do we know which PTE's to clear?



- Accessed bit set to 1 whenever accessed (done by hardware)

- Periodically the bit is set to 0
- Second chance algorithm - when the accessed bit is 0, it's given a second chance
- ↳ if it's still 0 a second time then we evict it.



- POTENTIAL → traverse the linked list and look for an accessed bit == 1
- ↳ if there is one that's 1 then set all the A-bits to 0 and continue
- ↳ if we ever wrap around and everything is still accessed then this is called the clock algorithm
- ↳ computer is too busy → you need more RAM

- ↳ if you ever find a PPN where all of its VPN's have accessed == 0

↳ evict the page

- In p7 we quit if we get a page fault and don't know how to handle it

- ↳ what we really want to do is try to handle it

- You can let the user try and handle it

- ↳ ex. they could mmap or smth

- Similarly, what if a process wants to get run or called period, call (ex. a clock or sleep function)

- ↳ introducing... signals!

kill(pid, signal);

- ↳ if it isn't called from within a handler then it will simply kill the process

kill() is sys call

↳ signal()

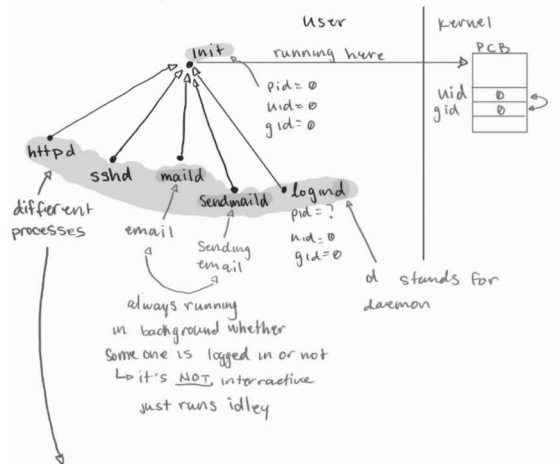
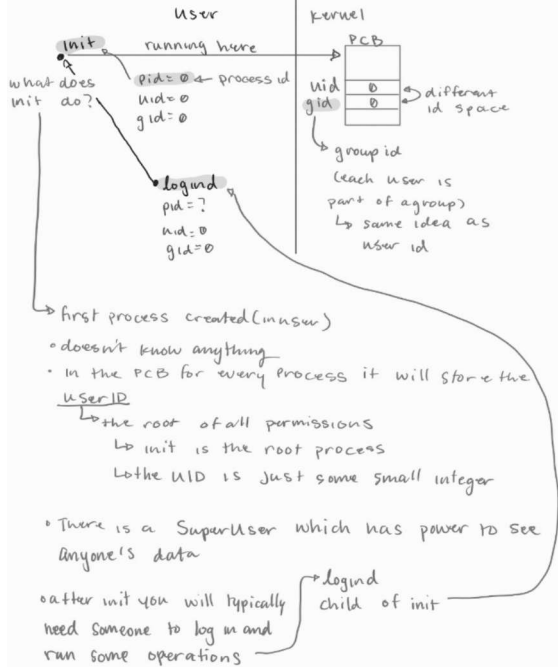
- Like an interrupt
- ↳ we can't trust the user to set up the IDT and stuff by itself so kernel does it for them

signal (signal, handler) diff nums for diff issues (segfault, illegal instruction, etc.)

void handler (int signal) run in user mode

↳ when done → signalreturn(), ← sys call: returns to failed instruction

if you don't call it will implicitly call signalreturn at the end



- how do these different processes get started?
- there is some sort of configuration folder with bunch of different binary files and stuff that tells how to start these processes
 - ↳ these files will also say how to restart in case of failures
 - ↳ they are also all formatted differently
- init starts a config daemon which finds the config files and starts the different config files

- login d
- given uid=0 which gives the process its super power to access sensitive files
 - ↳ it's CPL is still in user mode!

How do you authenticate + authorize users

- the user proves its identity
- we decide what powers this user has
- username + password
- ↳ proving you have information that ONLY that user will have
- ↳ what's the login process look like?

```
printf("User name: ");
scanf("%s", &s);
```

- this will tell the login daemon that someone is trying to login with a particular uid

```
printf("password: ");
scanf("%s", &s);
```

- Note that for username the terminal will echo your input + put your typed characters on the output...

BUT NOT for password now?

- ioctl special sys call that changes the terminal settings → can change the terminal to echo or not

- raw mode → no echo
- cooked mode → echo -- usually in this mode

How does the shell allow you to backspace?

- there is a character buffer that stores your inputs
 ↳ NOT in the actual file
- doesn't commit your characters until you press enter and commit the whole line
 ↳ this is why shell is line based
- other editors will encode the backspace as an actual ascii character and let the IDE handle it
- if the terminal has tab completion it's different
 ↳ has to be in raw mode b/c the shell can't handle the complex operation
- ↳ raw mode allows the program to handle the output by itself

/etc/passwd

- All users can read
 - Only root can write
 - one line per user
- IS: gheith: Ahmed Gheith: 30: password
- uid username realname pgid password
- internally, program only uses uid → username + real name are only for convenience
- now we only use a one-way hash

why isn't /etc/passwd only readable by root?

it's convenient for the process to have access to its uid + other user info

one way hash:
 easy to go from actual to encoded → Hard to go from encoded to actual

the hashing function is publicly known

- Supposedly the password input will read the password, hash it, and IMMEDIATELY forget it
- ↳ this WILL fail -- since we're using vmm
- ↳ when you page out the password is still there in physmem → can be accessed

Suppose you have an AMAZING hash function


- But someone on the system may choose a password from the dictionary
- Ex. they use the word "please"
- ↳ maybe you replace the l with a 1 or write it backwards... etc
- ↳ you could write a program that takes all the common possible values and apply the hash function
- ↳ then you could compare all the possible values to all the hashed passwords
- ↳ eventually you will get a hit
- ↳ once you get in you can use other methods to hack other things
- ↳ Solution: introduce randomization
- ↳ have some random value associated with each user
- ↳ how to use this dictionary method you have to iterate over everything again which makes it an unusable system again

All large systems will use LDAP

(or something like it)

- server that manages the users
- ↳ too many users to fit in one file

Let's switch back...



* now that we can log in we need to run some program

↳ how do we tell the system what program to run?

↳ in etc/passwd we store the default program

IS: gheith: Ahmed Gheith: 30: password: shell

chsh operation: change shell

* most of the time... the default shell is /bin/bash
↳ back in the day it was /ssh/shell

* After login, how does it create a new shell?

• it forks!

↳ parent: wait()

↳ when child quits, parent takes over and allows process to keep running

↳ child: runs shell

↳ exec("/bin/bash")
↳ shell

↳ when it enters shell it will be set up with root permissions

↳ logind uid=0

↳ before exec() it needs to set

things up to NOT be root

↳ change uid to whoever just logged in

↳ chdir to user home dir (also in /etc/passwd)

↳ after setup, switch to bash

IS: gheith: Ahmed Gheith: 30: password: shell: home

Now... What if we type ls?

* After the whole line is typed AND you press enter
↳ driver takes the line and gives it to the

shell program

↳ tokenizes by whitespace

• Assumes the first token is the command name

> ls -a abc

• tokenizes into

ls <-a> <abc>
↑
command name args

• There are a handful of built in command

↳ cd is built in

↳ ls is NOT built in

↳ MOST commands AREN'T built in

• If it is built in → just run it

• If it is NOT built in

↳ find a program with the command name

↳ exec to the program + pass the all the args

• how does it find the program?

↳ there is a subset of folders that it will check

↳ PATH=/bin/sbin/usr/bin

↳ some systems use colon instead of slash

↳ if you don't find the program → error

↳ then it exec to the program

↳ BUT if you exec then you lose the shell so...

fork() then exec()

↳ parent stays on shell and wait()

vfork()

• scary

• fork but don't copy address space...

↳ uses SAME stack as parent

• how do you prevent them from clashing?

↳ child "borrows" the address space

↳ prevents parent from running

↳ until child exec or exit

↳ child has to keep the stack clean

• use very rarely for very small operations

• now... COW is pretty good so we don't use

vfork() very much anymore.